

Group 96 [Raza Abbas, Alkesh Chaturvedi, Harsh Bhatia]

Problem Statement: Predict class of Iris plant using a decision tree ## 1. Import the data from Iris Dataset (Links to an external site.) (1 point) ## 2. Consider all columns as independent variables and assign to variable X except the last column and consider the last ## column as dependent variable and assign to variable y. (1 point) ## 3. Remove columns which don't help the problem statement. (1 point) ## 4. Compute some basic statistical details like percentile, mean, standard deviation of dataset (1 point) ## 5. Encode all the categorical columns into numeric (1 point) ## 6. Do Feature Scaling on Independent variables (2 points) ## 7. Split the data into train and test dataset (1 point) ## 8. Use sklearn library to train on train dataset on decision tree and predict on test dataset (3 points) ## 9. Compute the accuracy and precision. (2 points)

Importing the required libraries

```
In [1]: import pandas as pd
import numpy as np
import statistics
from urllib.request import urlretrieve
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
import seaborn as sns
```

1. Import the data from Iris Dataset (Links to an external site.) (1 point)

import the data from iris dataset & assign url of the file

```
In [2]: iris = "http://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
```

```
In [3]: urlretrieve(iris)
```

```
Out[3]: ('C:\\Users\\ealkcha\\AppData\\Local\\Temp\\tmpwm93a7c5',
<http.client.HTTPMessage at 0x2563c2b3310>)
```

Observation: saves the file locally

Read file into a dataframe and print its head

```
In [4]: df=pd.read_csv(iris,sep=',',names=["sepal_length", "sepal_width", "petal_length", "petal_width", "class"])
```

```
In [5]: df
```

```
Out[5]:
```

	sepal_length	sepal_width	petal_length	petal_width	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
...
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

150 rows × 5 columns

Observation: file is opened as data frame, first four columns are attributes and last column is class variable

```
In [6]: # display first 5 records with headers & shape
df.head(), df.shape
```

```
Out[6]: (   sepal_length  sepal_width  petal_length  petal_width    class
0          5.1          3.5          1.4          0.2  Iris-setosa
1          4.9          3.0          1.4          0.2  Iris-setosa
2          4.7          3.2          1.3          0.2  Iris-setosa
3          4.6          3.1          1.5          0.2  Iris-setosa
4          5.0          3.6          1.4          0.2  Iris-setosa,
(150, 5))
```

Observation: Displays the first 5 rows. Data has 150 rows and 5 columns

In [7]: *# check the data types of each attributes*
df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   sepal_length    150 non-null    float64
1   sepal_width     150 non-null    float64
2   petal_length    150 non-null    float64
3   petal_width     150 non-null    float64
4   class           150 non-null    object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

Observation: Displays the data type of each variable - attributes are float and class variable is qualitative data type; there are no null values in the data

2. Consider all columns as independent variables and assign to variable X except the last column and consider the last column as dependent variable and assign to variable y. (1 point)

Print the first 5 records to check Variable X & Y

In [8]: X=df.iloc[:,0:4]
X.head()

Out[8]:

	sepal_length	sepal_width	petal_length	petal_width
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

Observation: Variable X is assigned to the first four columns

```
In [9]: Y=df.iloc[:,-1:]
        Y.head()
```

```
Out[9]:
```

	class
0	Iris-setosa
1	Iris-setosa
2	Iris-setosa
3	Iris-setosa
4	Iris-setosa

Observation: Variable Y is assigned to the last columns

3. Remove columns which don't help the problem statement. (1 point)

```
In [10]: # To get a correlation we need to convert our categorical variable into numerical one
        df["class"] =df["class"].astype('category').cat.codes
```

```
In [11]: df["class"].dtype
```

```
Out[11]: dtype('int8')
```

Observation: Data type of the the Dependent variable changed from categorical to numeric

```
In [12]: # Check correlation among independent variables and also between dependent variable and each of the independent variables
        df.corr()
```


```
Out[12]:
```

	sepal_length	sepal_width	petal_length	petal_width	class
sepal_length	1.000000	-0.109369	0.871754	0.817954	0.782561
sepal_width	-0.109369	1.000000	-0.420516	-0.356544	-0.419446
petal_length	0.871754	-0.420516	1.000000	0.962757	0.949043
petal_width	0.817954	-0.356544	0.962757	1.000000	0.956464
class	0.782561	-0.419446	0.949043	0.956464	1.000000

Observation: Correlation table shows high correlation among sepal_length, petal_length & petal_width. These variables also have high correlation with dependent variable

```
In [13]: # Plot heat map to visualize correlation among the variables
sns.heatmap(df.corr(), cmap='coolwarm', annot=True)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x2563c4270a0>
```



Observation: Heat map provides visual display of correlation and reconfirms - high correlation among sepal_length, petal_length & petal_width. These variables also have high correlation with dependent variable

From correlation analysis, petal width has the highest correlation with the dependent variable among all highly correlated variables sepal_length, petal_length & petal_width. # we remove petal length and sepal length and keep petal width and sepal width in our model.

```
In [14]: df=df[["sepal_width", "petal_width", "class"]]
df.head()
```

```
Out[14]:
```

	sepal_width	petal_width	class
0	3.5	0.2	0
1	3.0	0.2	0
2	3.2	0.2	0
3	3.1	0.2	0
4	3.6	0.2	0

Observation: Displays data post removal of two variables - sepal_length and petal_length

4. Compute some basic statistical details like percentile, mean, standard deviation of dataset (1 point)

```
In [15]: stat = df.describe()
```

```
In [16]: stat
```


```
Out[16]:
```

	sepal_width	petal_width	class
count	150.000000	150.000000	150.000000
mean	3.054000	1.198667	1.000000
std	0.433594	0.763161	0.819232
min	2.000000	0.100000	0.000000
25%	2.800000	0.300000	0.000000
50%	3.000000	1.300000	1.000000
75%	3.300000	1.800000	2.000000
max	4.400000	2.500000	2.000000

Observation: From above statistics, we conclude the following: (1) Presence of outlier in sepal width as max entry is greater than $\text{mean} + 3 \times \text{std}$ (2) petal width seems to have slight skewness as seen from the mean and median values however there are no outliers

```
In [17]: sns.boxplot(x=df['petal_width'])
```


```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x2563c3bed30>
```



Observation: Box plot shows Petal Width data is left skewed but no outliers in the data

```
In [18]: sns.boxplot(x=df['sepal_width'])
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x2563cc92490>
```



Observation: Box plot indicates 4 minor outliers in sepal width data. We are not treating the outliers.

5. Encode all the categorical columns into numeric (1 point)

```
df["class"] = df["class"].astype('category').cat.codes
```

Above code was used earlier in the program to convert categorical variable ("class") into numeric. The variable is now numeric can be verified as below:

```
In [19]: # df["class"] = df["class"].astype('category').cat.codes <Command used earlier
         in the program to encode categorical column into numeric>
         df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   sepal_width  150 non-null    float64
1   petal_width  150 non-null    float64
2   class        150 non-null    int8
dtypes: float64(2), int8(1)
memory usage: 2.6 KB
```

Observation: As given above, the class variable is "int". this was categorical variable that we converted to "int" while finding the correlation among the variables. The command used for this conversion was: `df["class"] = df["class"].astype('category').cat.codes`

6. Do Feature Scaling on Independent variables (2 points)

```
In [20]: X=df.iloc[:,0:2]
X.head()
```

Out[20]:

	sepal_width	petal_width
0	3.5	0.2
1	3.0	0.2
2	3.2	0.2
3	3.1	0.2
4	3.6	0.2

Observation: Above detail shows that scales of two attributes sepal_width and petal width are significantly different

```
In [21]: sc = StandardScaler()
X = sc.fit_transform(X)
X[:10]
```

Out[21]: array([[1.03205722, -1.31297673],
[-0.1249576 , -1.31297673],
[0.33784833, -1.31297673],
[0.10644536, -1.31297673],
[1.26346019, -1.31297673],
[1.95766909, -1.05003079],
[0.80065426, -1.18150376],
[0.80065426, -1.31297673],
[-0.35636057, -1.31297673],
[0.10644536, -1.44444497]])

Observation: Two attributes are scaled as seen above

7. Split the data into train and test dataset (1 point)

lets divide 80% of the data into training set to train the model and rest into testing set; 80:20 ratio

```
In [22]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.20, random_state=22)
```

```
In [23]: X_train.shape
```

Out[23]: (120, 2)


```
In [24]: Y_train.shape
```

```
Out[24]: (120, 1)
```

```
In [25]: X_test.shape
```

```
Out[25]: (30, 2)
```

```
In [26]: Y_test.shape
```

```
Out[26]: (30, 1)
```

Observation: 150 records are split into train and test data in the ratio of 80:20 randomly

8. Use sklearn library to train on train dataset on decision tree and predict on test dataset (3 points)

```
In [27]: classifier = DecisionTreeClassifier()
```

```
In [28]: classifier.fit(X_train, Y_train)
```

```
Out[28]: DecisionTreeClassifier()
```

Observation: model is trained on the training data using the sklearn library

```
In [29]: Y_pred = classifier.predict(X_test)
```

Observation: prediction is done on the test data

9. Compute the accuracy and precision. (2 points)

```
In [30]: print(confusion_matrix(Y_test, Y_pred))
```

```
[[ 6  0  0]
 [ 0  8  2]
 [ 0  2 12]]
```

Observation: Confution matrix is displayed as above

```
In [31]: accuracy_score(Y_test, Y_pred)
```

```
Out[31]: 0.8666666666666667
```

Observation: Model accuracy is 86.6%

```
In [32]: print(classification_report(Y_test, Y_pred))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	6
Iris-versicolor	0.80	0.80	0.80	10
Iris-virginica	0.86	0.86	0.86	14
accuracy			0.87	30
macro avg	0.89	0.89	0.89	30
weighted avg	0.87	0.87	0.87	30

Observation: Precision, recall and f1-score are tabulated as above. All the parameters have values greater than 80% indicating reasonable performance for the model